

TRUSTPAY

AI-Powered Fraud Detection & Real-Time Transaction Monitoring System

A publication-grade engineering report presenting the design, mathematical normalization, ACID row-level locking architecture, and empirical machine learning benchmarks of the TrustPay enterprise banking platform. Combines Boyce-Codd Normal Form (BCNF) relational integrity, sub-2ms LightGBM fraud inference, and Groq Llama 3.3 GenAI forensic explanations.

4 Tiers

Decoupled Topology

BCNF

8 Relational Tables

99.98%

ML Accuracy / 100% Prec

1.31 ms

Mean Inference Latency

PROJECT CONTRIBUTORS & STUDENT CREDENTIALS

Raghav Gupta

20243226

B.Tech (3rd Year) • CSE

Rishabh Srivastava

20243236

B.Tech (3rd Year) • CSE

Rishabh Singh

20243235

B.Tech (3rd Year) • CSE

Prince Keshari

20243218

B.Tech (3rd Year) • CSE

14-Feature Behavioral Pipeline

Extracts velocity, Z-score dispersion, Haversine spatial jump speed, device UUIDs, and balance drain ratio in < 15 ms.

Dual-AI Explainability Layer

Combines LightGBM 0-100 risk scoring with Groq Llama 3.3 70B natural language rationales for compliance triage with fail-safe error holds.

ACID Stored Procedures & Triggers

Guarantees zero double-spending via row-level locks ([FOR UPDATE](#)) during dispute settlements, with SQL check triggers.

Empirical Cross-Dataset Benchmarks

Benchmarked across 100k project transactions (99.98%), ULB European Cards (284k tx, 99.951%), and IEEE-CIS real-world dataset (30k tx).


Database: Aiven Cloud MySQL 8 (InnoDB)

Backend & AI: Express 5 • FastAPI • LightGBM

Verification: Production Verified v1.0

1. Executive Summary & Distributed System Topology

TRUSTPAY is an enterprise financial technology platform engineered to process high-throughput financial transactions while simultaneously executing multi-tiered artificial intelligence risk analysis. The system architecture coordinates four decoupled tiers: a modern React 18 single-page application, an Express 5 API gateway, a FastAPI microservice combining LightGBM machine learning and Groq Llama 3.3 GenAI analysis, and a normalized MySQL 8 relational ledger with row-level locking.




System Layer	Core Technology	Operational Responsibility	Status
Tier 1: Client UI	React 18, Vite, React Router v6	Customer payment forms, live telemetry, and Administrative Sentinel compliance cockpit	Active
Tier 2: API Gateway	Node.js, Express 5, Axios	Stateless JWT Bearer RBAC, device fingerprinting, geocoding, balance pre-validation, ledger routing	Active
Tier 3: Sentinel AI	Python 3.13, FastAPI, LightGBM	14-feature behavioral pipeline, ML continuous probability, and Groq Llama 3.3 contextual reasoning	Active
Tier 4: Relational DB	MySQL 8 (InnoDB), Aiven Cloud	BCNF normalized schema (8 tables), row-level locks (FOR UPDATE), triggers, and stored procedures	Active

1.1 Architectural Report Table of Contents

Section 1: Executive Summary & Topology Section 2: Relational Database (ER Model & BCNF) Section 3: Master Entity Catalog & Data Dictionary Section 4: Concurrency Control & Stored Procedures Section 5: Core Banking API Gateway Architecture Section 6: Frontend Layer & 20-Point UI/UX Audit	Section 7: 7-Phase End-to-End Surveillance Pipeline Section 8: Sentinel AI Intelligence Architecture Section 9: Empirical ML Accuracy & Latency Benchmarks Section 10: Real-World Cross-Dataset Benchmarks (ULB & IEEE) Section 11: Cloud Deployment & Academic Project Defense Sign-Off
---	---

2. Normalized Relational Database Architecture

The TrustPay relational database engine operates on MySQL 8 with InnoDB tables enforcing strict foreign key constraints. The schema satisfies Boyce-Codd Normal Form (BCNF), eliminating redundancy, multi-valued dependencies, and transitive update anomalies across 8 distinct entities.



2.1 Relational Normalization Calculus (1NF to BCNF)

- **1NF (Atomic Attributes):** All table columns represent scalar, atomic values. Geolocation coordinates (lat/long) and client device fingerprints are decomposed into dedicated tables rather than serialized or delimited strings.
- **2NF (Full Functional Dependency):** All non-prime attributes are completely dependent on candidate keys. Composite natural keys were extracted into distinct entities with surrogate primary keys.
- **3NF (Transitive Dependency Removal):** Non-key attributes depend exclusively on primary keys ($X \rightarrow Y$ with no $Y \rightarrow Z$). Geolocation attributes were decoupled into `locations`, preventing city/country duplication within `transactions`.
- **BCNF (Strict Determinant Keys):** For every functional dependency $X \rightarrow Y$, X is a superkey. There are no overlapping candidate keys or dependencies on non-candidate attributes.

Referential Integrity Strategy: Child operational records (`accounts`, `devices`, `fraud_predictions`, `fraud_alerts`) cascade on user deletion. Conversely, `transactions` utilizes `ON DELETE RESTRICT` on accounts and `ON DELETE SET NULL` on devices and locations, guaranteeing audit survivability.

2.2 Master Entity Catalog & Data Dictionary

The relational ledger comprises 8 Boyce-Codd Normal Form (BCNF) compliant entities enforcing strict referential integrity, automated check triggers, and index optimization.

Entity Name	Primary Key	Foreign Keys & Cascades	Key Constraints	Business Role
users	user_id (BIGINT)	None	UNIQUE(email), UNIQUE(phone), role IN ('CUSTOMER', 'ADMIN')	Customer & Admin master identity
accounts	account_id (BIGINT)	user_id -> users (CASCADE)	UNIQUE(account_number), balance >= 0	Bank account liquidity ledger
devices	device_id (BIGINT)	user_id -> users (CASCADE)	UNIQUE(user_id, device_identifier)	Hardware profiles & client trust
locations	location_id (BIGINT)	None	Decoupled GPS lat/long & city attributes	Spatial master directory
transactions	transaction_id (BIGINT)	account_id (RESTRICT) device_id (SET NULL) location_id (SET NULL)	CHECK(amount > 0), status IN ('PENDING', 'APPROVED', 'REJECTED')	Core immutable financial ledger
fraud_predictions	prediction_id (BIGINT)	transaction_id (CASCADE)	CHECK(fraud_probability BETWEEN 0 AND 1), CHECK(risk_score BETWEEN 0 AND 100)	Output of LightGBM ML inference
fraud_alerts	alert_id (BIGINT)	transaction_id (CASCADE) prediction_id (CASCADE)	severity IN ('LOW', 'MEDIUM', 'HIGH', 'CRITICAL'), status IN ('OPEN', 'RESOLVED')	High-risk security triage queue
admin_reviews	review_id (BIGINT)	transaction_id (RESTRICT) admin_id -> users (RESTRICT)	decision IN ('APPROVED', 'REJECTED')	Immutable compliance audit log

2.3 Active Database Triggers

- trg_transaction_before_insert** (BEFORE INSERT ON **transactions**): Enforces monetary integrity by intercepting transactions where `amount ≤ 0` and terminating execution with `SIGNAL SQLSTATE '45000'`.
- trg_fraud_alert_resolved** (BEFORE UPDATE ON **fraud_alerts**): Automates lifecycle auditing by setting `resolved_at = CURRENT_TIMESTAMP` whenever an alert transitions to `'RESOLVED'`.
- trg_fraud_alert_before_insert** (BEFORE INSERT ON **fraud_alerts**): Defensive fallback ensuring that if an insert contains a `NULL` status, it is automatically assigned `'OPEN'`.

2.4 Cardinality & Relationship Matrix

The architecture models `users → accounts` (1:N), `users → devices` (1:N), `accounts → transactions` (1:N), and `locations → transactions` (1:N). Downstream security evaluation establishes 1:1 relationships between `transactions → fraud_predictions`, `fraud_predictions → fraud_alerts`, and `transactions → admin_reviews`.

3. Concurrency Control, Stored Procedures & API Gateway

3.1 ACID Stored Procedures with Row-Level Locking

Row-Level Locking Protection: `ApproveTransaction` and `RejectTransaction` execute `SELECT ... FROM transactions WHERE transaction_id = ? FOR UPDATE` within an explicit transaction block, eliminating race conditions when multiple administrators review transactions simultaneously.

Procedure Name	Parameters	Concurrency Logic & State Guarantees
<code>ApproveTransaction</code>	<code>p_tx_id, p_admin_id, p_reason</code>	Acquires row lock (<code>FOR UPDATE</code>), validates status is <code>PENDING</code> , updates status to <code>APPROVED</code> , updates account balance atomically (credits deposit / debits payment), inserts into <code>admin_reviews</code> , and marks <code>fraud_alerts</code> as <code>RESOLVED</code> .
<code>RejectTransaction</code>	<code>p_tx_id, p_admin_id, p_reason</code>	Acquires row lock (<code>FOR UPDATE</code>), validates status is <code>PENDING</code> , updates status to <code>REJECTED</code> , preserves customer balance untouched, inserts into <code>admin_reviews</code> , and marks open alerts as resolved.
<code>GetCustomerTransactions</code>	<code>p_user_id</code>	Returns customer-scoped chronological transaction ledger joined with account numbers.
<code>GetSuspiciousTransactions</code>	None	Pulls active fraud monitoring queue for transactions flagged as <code>HIGH_RISK</code> or severity <code>HIGH/CRITICAL</code> .

3.2 Core Banking API Gateway Architecture (Node.js & Express 5)

The API Gateway provides stateless authentication, role-based access control, request sanitation, and transactional orchestration with connection pooling (10 concurrent threads over TLS/SSL).

Method	Endpoint	Authorization	Functional Responsibility
<code>POST</code>	<code>/api/auth/register</code>	Public	Registers user, hashes password (Bcrypt), creates account
<code>POST</code>	<code>/api/auth/login</code>	Public	Validates credentials, binds device fingerprint, returns JWT
<code>POST</code>	<code>/api/transactions</code>	<code>CUSTOMER</code>	Initiates payment, pre-validates balance, queries ML, updates ledger
<code>GET</code>	<code>/api/transactions/admin/all</code>	<code>ADMIN</code>	Omnichannel ledger across all accounts with diagnostic filters
<code>POST</code>	<code>/api/transactions/admin/:id/approve</code>	<code>ADMIN</code>	Executes <code>ApproveTransaction</code> stored procedure with row lock
<code>POST</code>	<code>/api/transactions/admin/:id/reject</code>	<code>ADMIN</code>	Executes <code>RejectTransaction</code> stored procedure with row lock
<code>GET</code>	<code>/api/health</code>	Public	Liveness health probe returning database and gateway SLA status

4. Frontend Presentation Layer & UI/UX Audit Matrix

4.1 React 18 + Vite SPA Architecture

The presentation layer is engineered with React 18 and Vite for sub-second build times and hot reloading. State is decoupled across three global providers: `AuthContext` (stateless JWT, role RBAC, device UUID), `ThemeContext` (Light/Dark design tokens), and `ToastContext` (transactional alert feedback). Pure CSS design tokens eliminate runtime CSS-in-JS overhead while delivering an enterprise fintech aesthetic.


4.2 20-Point Fintech & Security Audit Matrix

A rigorous 20-point UI/UX audit was conducted across four standardized viewports (Desktop 1440px, Laptop 1280px, Tablet 768px, Mobile 390px) to verify visual ergonomics, accessibility, and error resilience.

#	Audit Criterion	Standard Specification	TrustPay Production Implementation	Result
1	Spacing Uniformity	Tokenized 4px/8px grid	CSS variables <code>--space-xs</code> to <code>--space-2xl</code> enforce uniform margins	Passed
2	Typography Scale	Inter / Plus Jakarta Sans	6-level type hierarchy with strict line-height and weight tracking	Passed
3	Card Sizing	Responsive aspect ratios	Auto-fit CSS grid templates prevent card clipping across breakpoints	Passed
4	Sidebar Behavior	Persistent / Drawer	Fixed desktop navigation with smooth collapsible backdrop drawer on mobile	Passed
5	Responsive Tables	Horizontal scroll safety	Custom wrappers with touch affordance, sticky headers, and mobile cards	Passed
6	Button States	Clear interactive states	Standardized primary/secondary buttons with <code>:focus-visible</code> rings	Passed
7	Loading States	Zero cumulative layout shift	Dual-ring SVG spinners and pulse skeleton placeholders for async calls	Passed
8	Modal Dialogs	Clamped viewport bounds	Centered modals with backdrop blur (<code>blur(8px)</code>) and escape key listeners	Passed
9	Accessibility	WCAG 2.1 AA Compliance	ARIA labels on icon triggers, semantic headings, keyboard focusable controls	Passed
10	Contrast Ratios	Minimum 4.5:1 ratio	Verified contrast compliant in both Dark Mode (<code>#0b0f19</code>) and Light Mode	Passed
11	Fraud Risk Meter	Intuitive visual telemetry	Calibrated 0–100 risk score meter with color gradient and deviation badges	Passed
12	Admin Workflow	Step-by-step triage	Unified dossier: AI reasoning, phone dialer simulation, mandatory audit reason	Passed
13	Mobile Ergonomics	Touch target ≥ 44px	Zero horizontal body scroll, thumb-accessible primary buttons on mobile	Passed

5. Complete 7-Phase Transaction Surveillance Pipeline

TrustPay coordinates a 7-phase execution pipeline spanning client inputs, database constraints, behavioral feature pipelines, machine learning inference, generative AI explanations, and administrative dispute settlements.




5.1 Phase-by-Phase Execution Breakdown

- Phase 1: Client Submission & Telemetry Ingestion** — Retail customer submits transaction. HTML5 Geolocation API captures client coordinates (lat/long) and local storage supplies immutable device fingerprint UUID.
- Phase 2: Gateway Ingestion & Pre-Validation** — Express gateway verifies Bearer JWT and extracts user ID. For debit operations, gateway asserts `balance ≥ amount`; if insufficient, HTTP 400 is returned immediately.
- Phase 3: Ledger Hold & Trigger Validation** — Transaction record inserted with status `'PENDING'`. Database trigger `trg_transaction_before_insert` validates `amount > 0`.
- Phase 4: Dual-AI Intelligence Inference** — FastAPI microservice queries customer history, computes 14 behavioral features, runs LightGBM classifier for continuous probability and 0-100 score, and invokes Groq Llama 3.3 for forensic explanations.
- Phase 5: Hybrid Decision Branching:**
 - *Branch A (Low/Med Risk, Score < 70):* Auto-approved, balance debited/credited, HTTP 201 instant settlement.
 - *Branch B (High Risk, Score ≥ 70):* Security alert inserted in `fraud_alerts`, status stays `PENDING`, funds preserved.
 - *Branch C (Fail-Safe Error Hold):* Fallback if ML service times out. Status stays `PENDING_ERROR`, funds safe.
- Phase 6: Compliance Audit & Row-Locked Stored Procedure** — Compliance officer inspects anomaly in `/admin/transactions`, enters mandatory justification, and calls stored procedure `ApproveTransaction` or `RejectTransaction` with row lock (`FOR UPDATE`).
- Phase 7: Real-Time Reconciliation** — Balance settlement finalized, alert timestamped as resolved by trigger, and dashboards updated across platform.

6. Sentinel AI Intelligence Layer & Feature Pipeline

The fraud intelligence layer combines mathematical feature extraction, a gradient-boostered decision tree ensemble, and large language model reasoning to provide explainable fraud decisions.



6.1 14-Feature Behavioral Engineering Catalog

Feature Name	Type	Mathematical Formulation	Fraud Detection Signal
amount	Continuous	Raw transaction magnitude (X)	Monetary exposure
customer_mean_amount	Continuous	$\mu = (1/N) \sum X_i$	Historical spending baseline
customer_std_amount	Continuous	$\sigma = \sqrt{(1/N) \sum (X_i - \mu)^2}$	Spending volatility
amount_to_mean_ratio	Continuous	$R = X / \max(\mu, 1.0)$	Relative spend spike factor
is_unusual_amount	Binary (0/1)	1 if $X > \mu + 2.5\sigma$ or $R > 4.0$, else 0	Extreme monetary anomaly flag
hour_of_day / is_unusual_time	Discrete/Bin	$H \in [0, 23]$; 1 if H outside user active band	Temporal and off-hours anomaly
known_device / new_device	Binary (0/1)	1 if DeviceID $\in$ Devices_user, else 0	Account takeover (ATO) indicator
known_location / location_changed	Binary (0/1)	1 if Distance(Loc_curr, Loc_prev) > 300 km	Impossible travel / proxy jumping
transaction_count	Integer	Cumulative lifetime transactions (N)	Account maturity index
history_confidence	Categorical	NONE (N=0), NEW ($1 \leq N < 5$), ESTABLISHED	Cold-start confidence tiering

6.2 Deterministic Decision Engine Rules (engine.py)

The final transaction settlement action fuses calibrated ML probabilities with deterministic safety rules:

- **Cold-Start Safety Guard:** First-time accounts (N=0) with score ≥ 70 are routed to `ADMIN_REVIEW`; otherwise placed under `MONITOR`.
- **Impossible Travel Teleportation:** Real-world speed $V > 800$ km/h over distance $\Delta d > 100$ km floors risk score at 88.0 and forces `ADMIN_REVIEW`.
- **ATO Triad:** Unrecognized device combined with location shift and balance drain ratio $> 85\%$ triggers mandatory administrative security hold.
- **Rapid Burst Testing:** ≥ 3 transactions in 10 minutes from a new device triggers security flag.
- **Familiar Baseline Instant Clearance:** Known device, recognized location, normal spend ratio, and ML score < 30 auto-clears instantly.

7. Empirical Machine Learning & Decision Layer Benchmarks

The production ML engine was evaluated across 100,000 transactions using a strict 80/20 chronological holdout test set (20,000 unseen transactions: 19,311 clean, 689 fraud; 28:1 class imbalance) measured via hardware timers (`time.perf_counter_ns`).

7.1 Statistical Accuracy Metrics (Holdout Test Set N = 20,000)

Evaluation Metric	Measured Score	Industry Standard	Operational Significance
Overall Accuracy	99.98%	> 98.0%	Near-zero global misclassification rate
Balanced Accuracy	99.64%	> 95.0%	High discrimination across skewed class distributions
ROC-AUC Score	0.9991	> 0.950	High class separability between fraud and clean
PR-AUC (Avg Precision)	0.9986	> 0.900	Precision sustained across all sensitivity thresholds
Fraud Precision	100.00%	> 92.0%	Zero false alarms ; legitimate transactions never flagged
Fraud Recall (Sensitivity)	99.27%	> 90.0%	684 of 689 fraud attempts intercepted directly by model
Decision Layer Total Intercept	99.42%	> 92.0%	685 of 689 fraud attempts caught across Monitor + Admin Review
Clean Auto-Approval Rate	98.96%	> 95.0%	19,111 legitimate users cleared instantly without friction
Admin Hard-Block Precision	100.00%	> 98.0%	Zero legitimate customers ever falsely hard-blocked

7.2 Multi-Tier Latency Benchmarks (ML Model + LLM + Decision Layer)

Pipeline Layer / Subsystem	Engine	Mean Latency	Median (p50)	Tail (p99)	Throughput
1. Feature Extraction	NumPy / Pandas	0.127 ms	0.123 ms	0.192 ms	7,874 QPS
2. ML Model + Physics Bounds	LightGBM 300t	2.346 ms	2.297 ms	3.610 ms	426 QPS
• Raw Tree (<i>predict_proba</i>)	C++ Booster	1.310 ms	1.200 ms	1.840 ms	764 QPS
3. GenAI Context Builder	JSON Serializer	0.036 ms	0.035 ms	0.058 ms	27,770 QPS
4. LLM Forensic Analyst	Groq Llama 3.3 70B	1,435.6 ms	1,452.2 ms	1,541.8 ms	~0.7 QPS
• Zero-Delay Fallback	Deterministic Rule	0.005 ms	0.005 ms	0.009 ms	200,000 QPS
5. Deterministic Decision Engine	Institutional Matrix	0.008 ms	0.008 ms	0.014 ms	125,000 QPS
Fast-Path Settlement (Steps 1+2+5)	Core Local Stack	2.517 ms	2.465 ms	3.859 ms	561.4 QPS
Full Forensic Pipeline (Steps 1..5)	ML + Groq LLM	1,438.1 ms	1,454.7 ms	1,545.7 ms	Async / Audit
Batch Offline Audit (20k tx)	LightGBM Matrix	3.50 µs / item	—	—	285,318 QPS

7.3 Decision Threshold Sensitivity & Feature Importance

Confusion Matrix (Holdout N = 20,000) - Actual Clean (19,311): 19,311 True Negatives (100.0%), 0 False Positives - Actual Fraud (689): 684 True Positives (99.27%), 5 False Negatives	Top Feature Split Gains: <code>amount</code> (1,801 splits) <code>amount_to_historical_mean</code> (1,305 splits) <code>amount_zscore</code> (995 splits) <code>transaction_count</code> (877 splits)
--	--

8. Real-World Cross-Dataset Generalization Benchmarks

To evaluate portability across diverse financial distributions, the TrustPay LightGBM architecture was benchmarked against two public real-world datasets: the **ULB European Cardholders dataset** (284,807 transactions) and the **IEEE-CIS Fraud Detection dataset** (30,000 transactions).

Evaluation Dimension	TrustPay Sentinel (Project Baseline)	ULB European Credit Cards (Kaggle)	IEEE-CIS Dataset (Full 48 Features)
Dataset Domain	Simulated Core Banking	Real-World Credit Cards	Real-World E-Commerce & CNP
Total Transactions	100,000	284,807	30,000
Class Imbalance	3.48% (1:28)	0.173% (1:578)	2.87% (1:34)
Features Evaluated	20 Behavioral & Telemetry	30 (PCA Components)	48 (Card, Addr, Vesta)
Overall Accuracy	99.98%	99.951%	97.97%
ROC-AUC Score	0.9991	0.9770	0.8994
PR-AUC Score	0.9986	0.7651	0.6113
Fraud Precision	100.00%	91.23%	90.00%
Operational Recall	99.42% (Decision Layer)	74.67%	68.82%
Clean Specificity	100.00%	99.991%	99.86%
Mean Latency	1.31 ms	1.58 ms	1.90 ms
Batch QPS	285,318 QPS	366,069 QPS	183,940 QPS

8.1 Comparative Evaluation with 20 Strict Project Features

When the IEEE-CIS dataset was stripped of its native attributes and evaluated **strictly on TrustPay's 20 engineered telemetry features**, the model attained **96.90% Accuracy**, **0.7575 ROC-AUC**, **60.00% Precision** (at $\tau=0.30$), **100.00% Specificity**, and **0.85 ms inference latency**.

8.2 Architectural Analysis of Generalization

Why TrustPay Achieves 99.98% Accuracy: In modern core banking, the platform possesses direct, unblinded client telemetry: persistent hardware device UUIDs, HTML5 GPS coordinates, customer account balance drain ratios, and dedicated customer spend baselines. This rich context enables the hybrid LightGBM + Decision Engine architecture to separate legitimate behavior from unauthorized attacks with **zero false alarms** (100% precision) and **99.42% operational fraud capture**.

Generalization on Blinded Real-World Vectors (ULB Dataset): On the ULB benchmark, all bank telemetry was privacy-anonymized into mathematical PCA components (V1 to V28). Even under an extreme **0.173% class imbalance (1 fraud in 578 clean)**, the TrustPay LightGBM architecture achieved **0.9770 ROC-AUC**, **99.951% accuracy**, and **91.23% precision** (only 5 false alarms across 56,887 legitimate transactions).

9. Cloud Production Deployment & Academic Project Defense Sign-Off

9.1 Cloud Infrastructure Deployment Topology

Infrastructure Tier	Cloud Provider & Host	Configuration & Protocol	SLA & Security
Core API Gateway	Render Cloud Platform	Node.js 20 LTS, Express 5, HTTPS/REST	TLS 1.3, Automated Health
Relational Database	Aiven Cloud	MySQL 8.0.35 Enterprise (InnoDB)	SSL Encrypted, PITR Backup
Sentinel AI Service	FastAPI Microservice	Python 3.13, Uvicorn ASGI, LightGBM	Sub-2ms Inference
SPA Client Hosting	Netlify / Vercel Edge	Static React 18 SPA with rewrite rules	Zero 404 Deep-Linking

9.2 Verification, Quality & Operational Matrix

Verification Domain	Test Procedure	Observed Result	Status
Frontend Quality	<code>npm run lint</code>	0 errors, 0 warnings	Passed
Frontend Production	<code>npm run build</code>	Vite build in 257ms	Passed
Backend Syntax	<code>node -c src/server.js</code>	Zero syntax errors	Passed
Database Triggers	Negative amount insertion harness	Blocked (SQLSTATE 45000)	Passed
Concurrency Locks	Concurrent approval stored procs	Row lock (FOR UPDATE) pass	Passed
Fail-Safe Fallback	Simulated ML timeout degradation	Safe hold in PENDING_ERROR	Passed

9.3 Academic Project Defense & Evaluation Sign-Off

This academic project report and software artifact are submitted in partial fulfillment of the requirements for **Bachelor of Technology (B.Tech 3rd Year) in Computer Science & Engineering**.

PROJECT GUIDE	HEAD OF DEPARTMENT	EXTERNAL EXAMINER
----- Signature & Date Department of CSE	----- Signature & Date Department of CSE	----- Signature & Date 3rd Year Project Evaluation Board